

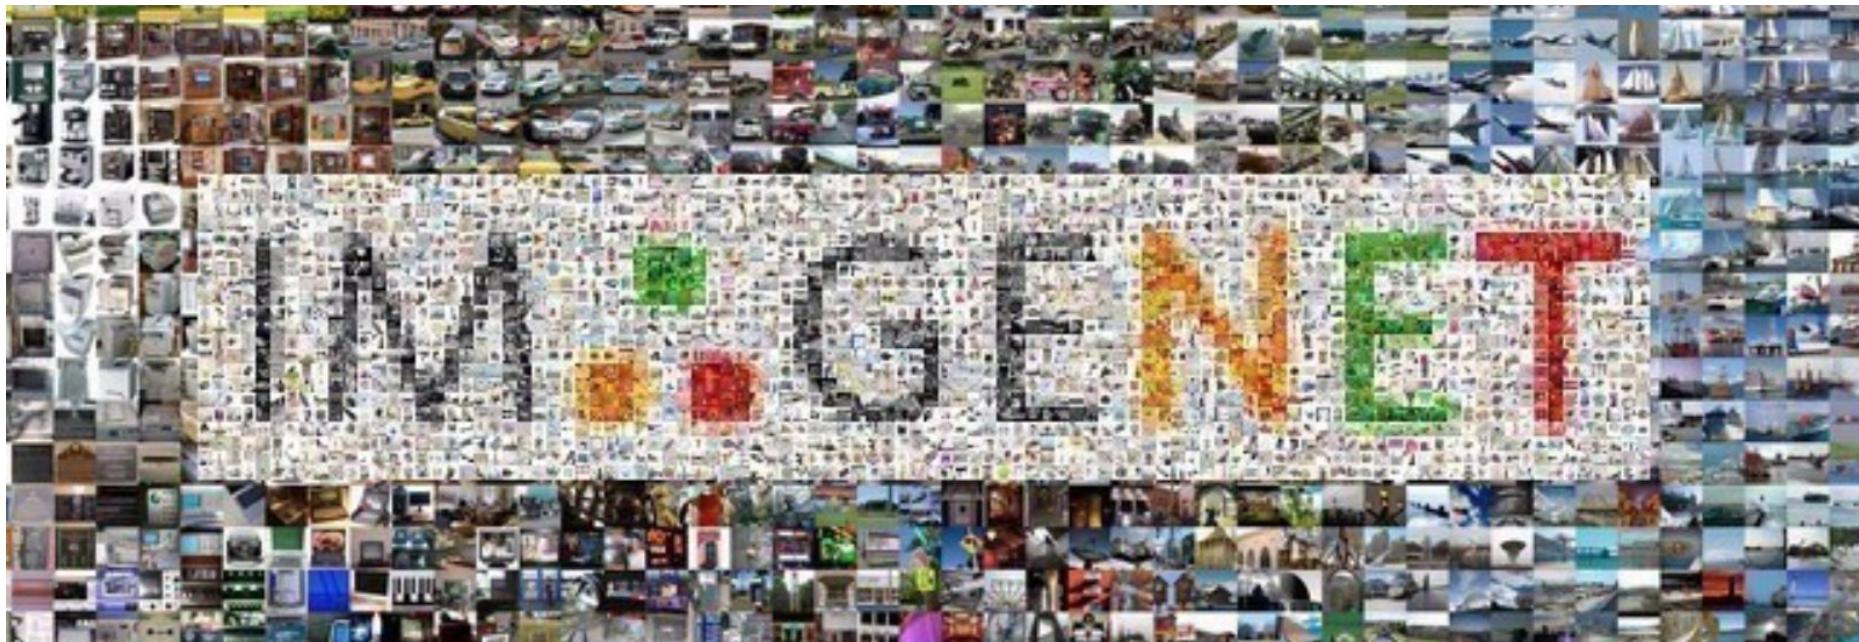
# META-LEARNING FOR SEMI-SUPERVISED FEW-SHOT CLASSIFICATION

**Mengye Ren<sup>†✉</sup>, Eleni Triantafillou<sup>\*†✉</sup>, Sachin Ravi<sup>\*§</sup>, Jake Snell<sup>†✉</sup>, Kevin Swersky<sup>¶</sup>,  
Joshua B. Tenenbaum<sup>¤</sup>, Hugo Larochelle<sup>¶‡</sup> & Richard S. Zemel<sup>†‡✉</sup>**

<sup>†</sup>University of Toronto, <sup>§</sup>Princeton University, <sup>¶</sup>Google Brain, <sup>¤</sup>MIT, <sup>‡</sup>CIFAR, <sup>✉</sup>Vector Institute  
`{mren,eleni}@cs.toronto.edu, sachinr@cs.princeton.edu,`  
`jsnell@cs.toronto.edu, kswersky@google.com,`  
`jbt@mit.edu, hugolarochelle@google.com, zemel@cs.toronto.edu`

CS330 Paper Presentation: October 16th, 2019

# Supervised Classification

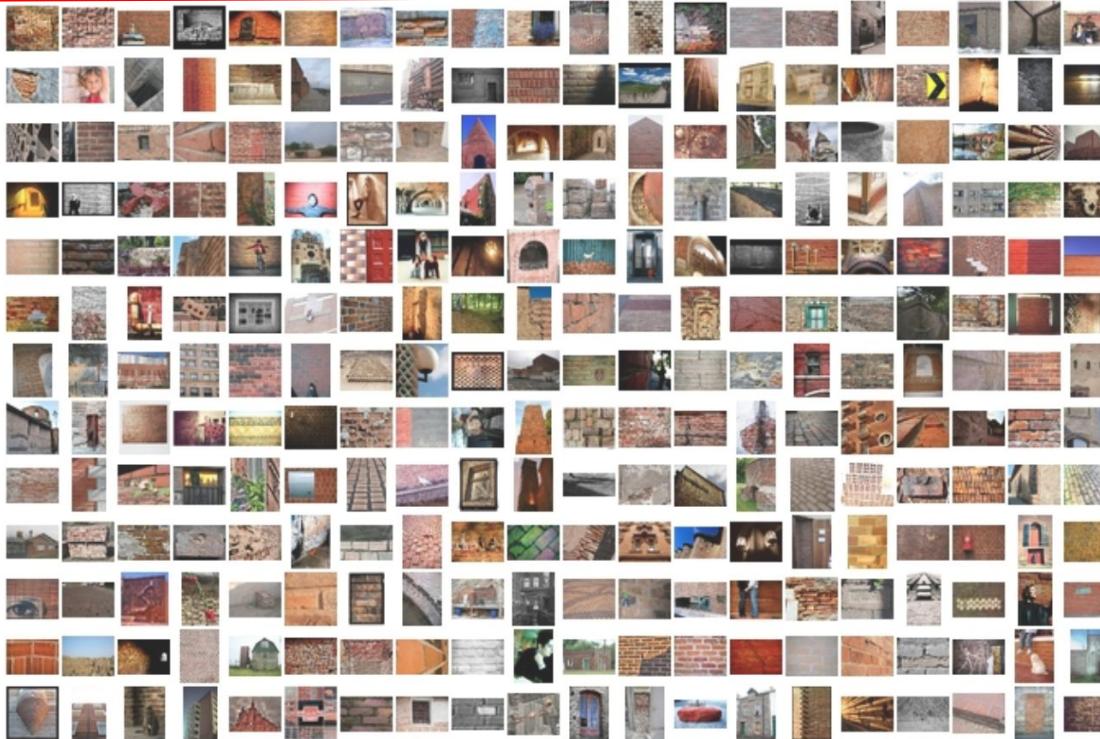


# Semi-Supervised Classification: More realistic dataset

Labelled



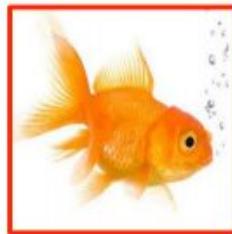
Unlabelled



# Semi-Supervised Classification

Most “biologically plausible” learning regime

# A familiar problem:



“goldfish”



“shark”

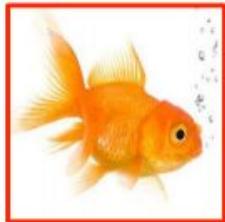
Support Set

**Few-shot, multi-task learning:**  
Generalize to unseen classes

?



# A new twist on a familiar problem:



“goldfish”



“shark”



Unlabeled Set

?



How can we leverage unlabelled data  
for few-shot classification?

**Training**

1



2



3



Support Set

+



+



-



+



-



?



?



Query Set

**Testing**

1



2



3



Support Set

-



+



-



+



+



?



?



?



Query Set

Unlabelled data may come from the support set or not  
**(distractors)**

# Strategy:

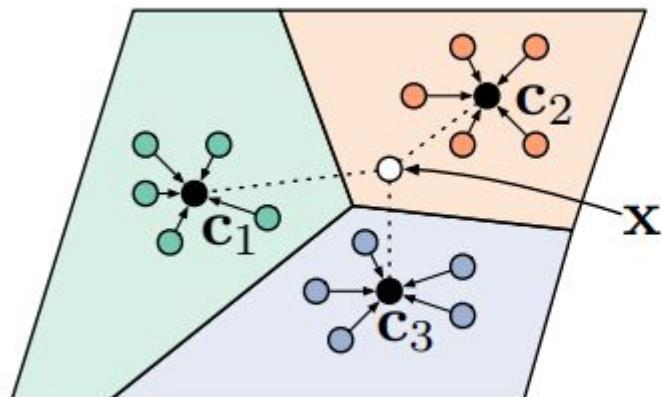
As we can now appreciate, there are a number of possible ways to approach the original problem. To name a few:

- Siamese Networks (Koch et al, 2015)
- Matching Networks (Vinyals et al., 2016)
- Prototypical Networks (Snell et al., 2017)
- Weight initialization / Update step learning (Ravi et al., 2017, Finn et al., 2017)
- MANN (Santoro et al., 2016)
- Temporal convolutions (Mishra et al., 2017)

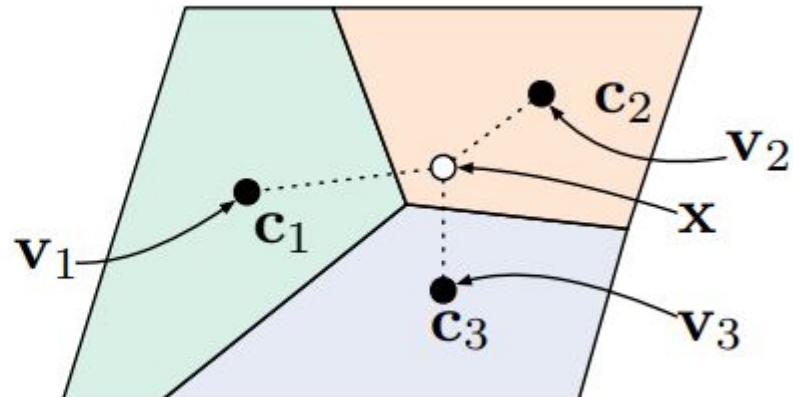
**All are reasonable starting points for semi-supervised few-shot classification problem!**

# Prototypical Networks (Snell et al., 2017)

Very simple inductive bias!



(a) Few-shot

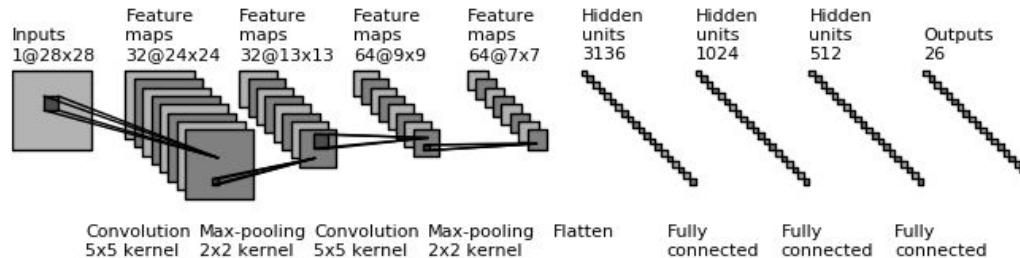


(b) Zero-shot

# Prototypical Networks (Snell et al., 2017)

$$\mathbf{c}_k = \frac{1}{|S_k|} \sum_{(\mathbf{x}_i, y_i) \in S_k} f_\phi(\mathbf{x}_i)$$

For each class, compute **prototype**



Embedding is generated via a simple convnet:

**Pixels - 64 [3x3] Filters - Batchnorm - ReLU - [2x2] MaxPool = 64D Vector**

# Prototypical Networks (Snell et al., 2017)

$$\mathbf{c}_k = \frac{1}{|S_k|} \sum_{(\mathbf{x}_i, y_i) \in S_k} f_\phi(\mathbf{x}_i)$$

For each class, compute **prototype**

$$p_\phi(y = k \mid \mathbf{x}) = \frac{\exp(-d(f_\phi(\mathbf{x}), \mathbf{c}_k))}{\sum_{k'} \exp(-d(f_\phi(\mathbf{x}), \mathbf{c}_{k'}))}$$

Softmax distribution of distances to **prototypes** for new image

$$J(\phi) = -\log p_\phi(y = k \mid \mathbf{x})$$

Compute loss

# Prototypical Networks (Snell et al., 2017)

$$\mathbf{c}_k = \frac{1}{|S_k|} \sum_{(\mathbf{x}_i, y_i) \in S_k} f_\phi(\mathbf{x}_i)$$

For each class, compute **prototype**

$$p_\phi(y = k \mid \mathbf{x}) = \frac{\exp(-d(f_\phi(\mathbf{x}), \mathbf{c}_k))}{\sum_{k'} \exp(-d(f_\phi(\mathbf{x}), \mathbf{c}_{k'}))}$$

Softmax distribution of distances to **prototypes** for new image

$$J(\phi) = -\log p_\phi(y = k \mid \mathbf{x})$$

Compute loss

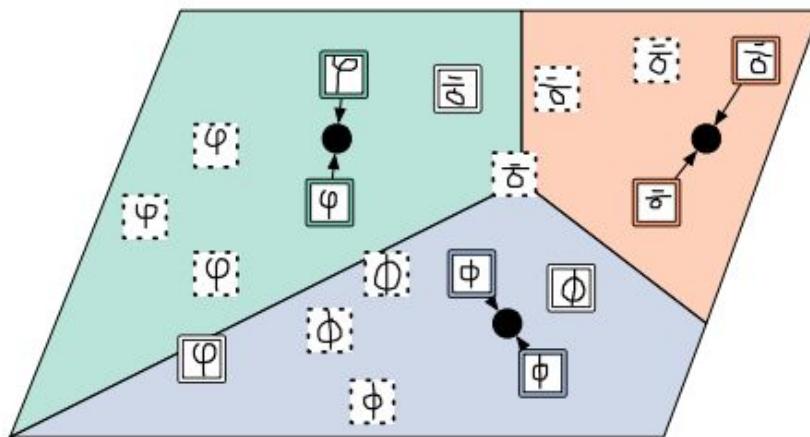
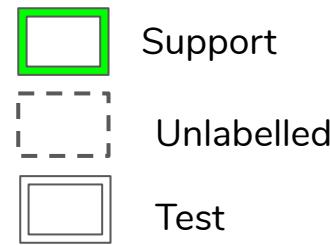
Very simple inductive bias:

Reduces to a linear model with Euclidean distance  $d(\mathbf{z}, \mathbf{z}') = \|\mathbf{z} - \mathbf{z}'\|^2$

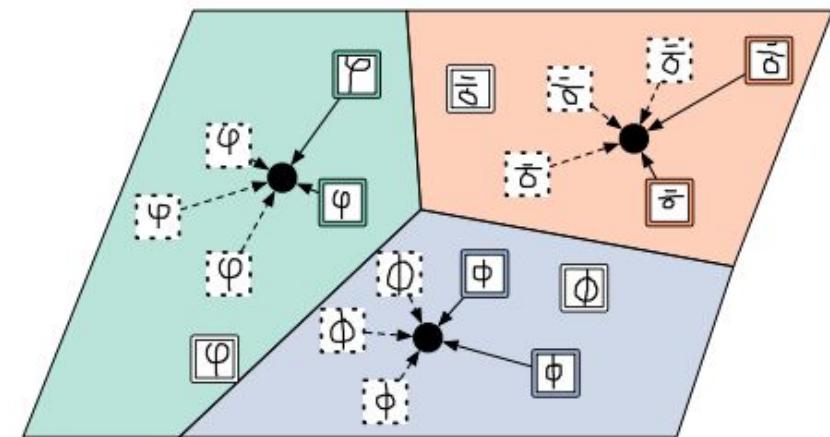
$$\mathbf{w}_k^\top f_\phi(\mathbf{x}) + b_k, \text{ where } \mathbf{w}_k = 2\mathbf{c}_k \text{ and } b_k = -\mathbf{c}_k^\top \mathbf{c}_k$$

# Strategy for semi-supervised:

Refine Prototypes centers with unlabelled data.



Before Refinement



After Refinement

# Strategy for semi-supervised:

1. Start with labelled prototypes
2. Give each unlabelled input a **partial assignment** to each cluster
3. Incorporate unlabelled examples into original prototype

# Prototypical networks with Soft k-means

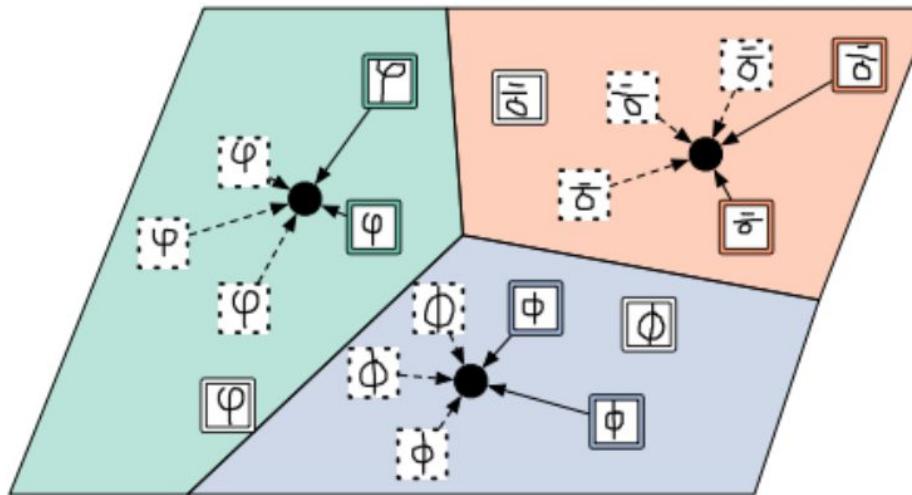
$$\mathcal{R} = \{\tilde{\mathbf{x}}_1, \tilde{\mathbf{x}}_2, \dots, \tilde{\mathbf{x}}_M\}$$

Unlabelled support set

$$\tilde{\mathbf{p}}_c = \frac{\sum_i h(\mathbf{x}_i) z_{i,c} + \sum_j h(\tilde{\mathbf{x}}_j) \tilde{z}_{j,c}}{\sum_i z_{i,c} + \sum_j \tilde{z}_{j,c}}, \text{ where } \tilde{z}_{j,c} = \frac{\exp(-\|\mathbf{h}(\tilde{\mathbf{x}}_j) - \mathbf{p}_c\|_2^2)}{\sum_{c'} \exp(-\|\mathbf{h}(\tilde{\mathbf{x}}_j) - \mathbf{p}_{c'}\|_2^2)}$$

Partial Assignment

# Prototypical networks with Soft k-means



After Refinement

What about distractor classes?

# Prototypical networks with Soft k-means w/ Distractor Cluster

Add a **buffering prototype** at the origin to “capture the distractors”

$$\mathbf{p}_c = \begin{cases} \frac{\sum_i h(\mathbf{x}_i) z_{i,c}}{\sum_i z_{i,c}} & \text{for } c = 1 \dots N \\ \mathbf{0} & \text{for } c = N + 1 \end{cases}$$

# Prototypical networks with Soft k-means w/ Distractor Cluster

Add a **buffering prototype** at the origin to “capture the distractors”

$$\mathbf{p}_c = \begin{cases} \frac{\sum_i h(\mathbf{x}_i) z_{i,c}}{\sum_i z_{i,c}} & \text{for } c = 1 \dots N \\ \mathbf{0} & \text{for } c = N + 1 \end{cases}$$

Assumption: Distractors all come from one class!

# Soft k-means + Masking Network

1. Distance

$$\tilde{d}_{j,c} = \frac{d_{j,c}}{\frac{1}{M} \sum_j d_{j,c}}, \text{ where } d_{j,c} = \|h(\tilde{\mathbf{x}}_j) - \mathbf{p}_c\|_2^2$$

2. Compute mask with small network

$$[\beta_c, \gamma_c] = \text{MLP} \left( \left[ \min_j(\tilde{d}_{j,c}), \max_j(\tilde{d}_{j,c}), \text{var}_j(\tilde{d}_{j,c}), \text{skew}_j(\tilde{d}_{j,c}), \text{kurt}_j(\tilde{d}_{j,c}) \right] \right)$$

$$\tilde{\mathbf{p}}_c = \frac{\sum_i h(\mathbf{x}_i) z_{i,c} + \sum_j h(\tilde{\mathbf{x}}_j) \tilde{z}_{j,c} m_{j,c}}{\sum_i z_{i,c} + \sum_j \tilde{z}_{j,c} m_{j,c}}, \text{ where } m_{j,c} = \sigma \left( -\gamma_c (\tilde{d}_{j,c} - \beta_c) \right)$$

# Soft k-means + Masking Network

$$\tilde{d}_{j,c} = \frac{d_{j,c}}{\frac{1}{M} \sum_j d_{j,c}}, \text{ where } d_{j,c} = \|h(\tilde{\mathbf{x}}_j) - \mathbf{p}_c\|_2^2$$

$$[\beta_c, \gamma_c] = \text{MLP} \left( \left[ \min_j(\tilde{d}_{j,c}), \max_j(\tilde{d}_{j,c}), \text{var}_j(\tilde{d}_{j,c}), \text{skew}_j(\tilde{d}_{j,c}), \text{kurt}_j(\tilde{d}_{j,c}) \right] \right)$$

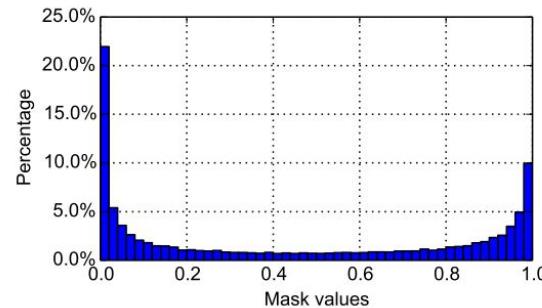
differentiable

$$\tilde{\mathbf{p}}_c = \frac{\sum_i h(\mathbf{x}_i) z_{i,c} + \sum_j h(\tilde{\mathbf{x}}_j) \tilde{z}_{j,c} m_{j,c}}{\sum_i z_{i,c} + \sum_j \tilde{z}_{j,c} m_{j,c}}, \text{ where } m_{j,c} = \sigma \left( -\gamma_c (\tilde{d}_{j,c} - \beta_c) \right)$$

# Soft k-means + Masking

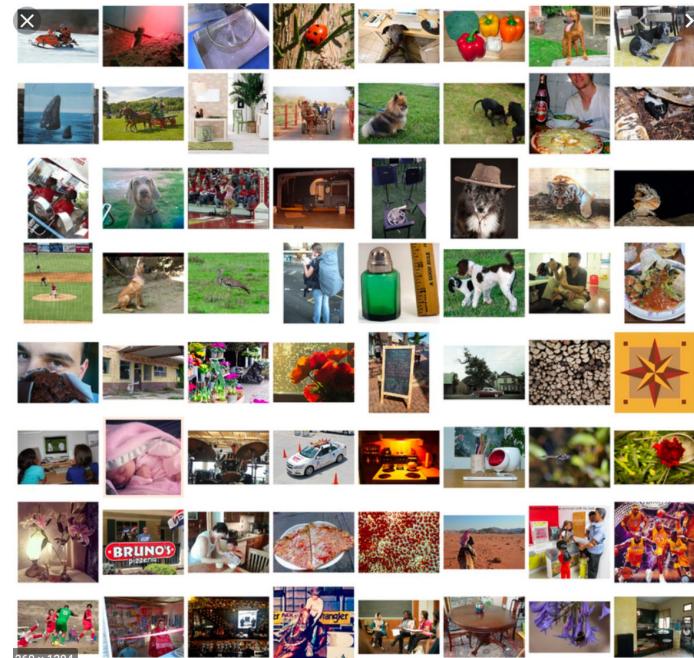
$$[\beta_c, \gamma_c] = \text{MLP} \left( \left[ \min_j(\tilde{d}_{j,c}), \max_j(\tilde{d}_{j,c}), \text{var}_j(\tilde{d}_{j,c}), \text{skew}_j(\tilde{d}_{j,c}), \text{kurt}_j(\tilde{d}_{j,c}) \right] \right)$$

In practice, MLP is a dense layer with 20 hidden units (tanh nonlinearity)



# Datasets

- Omniglot
  - minilmageNet (600 images from 100 classes)



# Hierarchical Datasets

Omniglot

**Train Alphabets:** Alphabet\_of\_the\_Magi, Angelic, Anglo-Saxon\_Futhorc, Arcadian, Asomtavruli\_(Georgian), Atemayar\_Qelisayer, Atlantean, Aurek-Besh, Avesta, Balinese, Blackfoot\_(Canadian\_Aboriginal\_Syllabics), Braille, Burmese\_(Myanmar), Cyrillic, Futurama, Ge\_ez, Glagolitic, Grantha, Greek, Gujarati, Gurmukhi (character 01-41), Inuktitut\_(Canadian\_Aboriginal\_Syllabics), Japanese\_(hiragana), Japanese\_(katakana), Korean, Latin, Malay\_(Jawi\_-\_Arabic), N\_Ko, Ojibwe\_(Canadian\_Aboriginal\_Syllabics), Sanskrit, Syriac\_(Estrangelo), Tagalog, Tifinagh

**Validation Alphabets:** Armenian, Bengali, Early\_Aramaic, Hebrew, Mkhedruli\_(Geogian)

**Test Alphabets:** Gurmukhi (character 42-45), Kannada, Keble, Malayalam, Manipuri, Mongolian, Old\_Church\_Slavonic\_(Cyrillic), Oriya, Sylheti, Syriac\_(Serto), Tengwar, Tibetan, ULOG

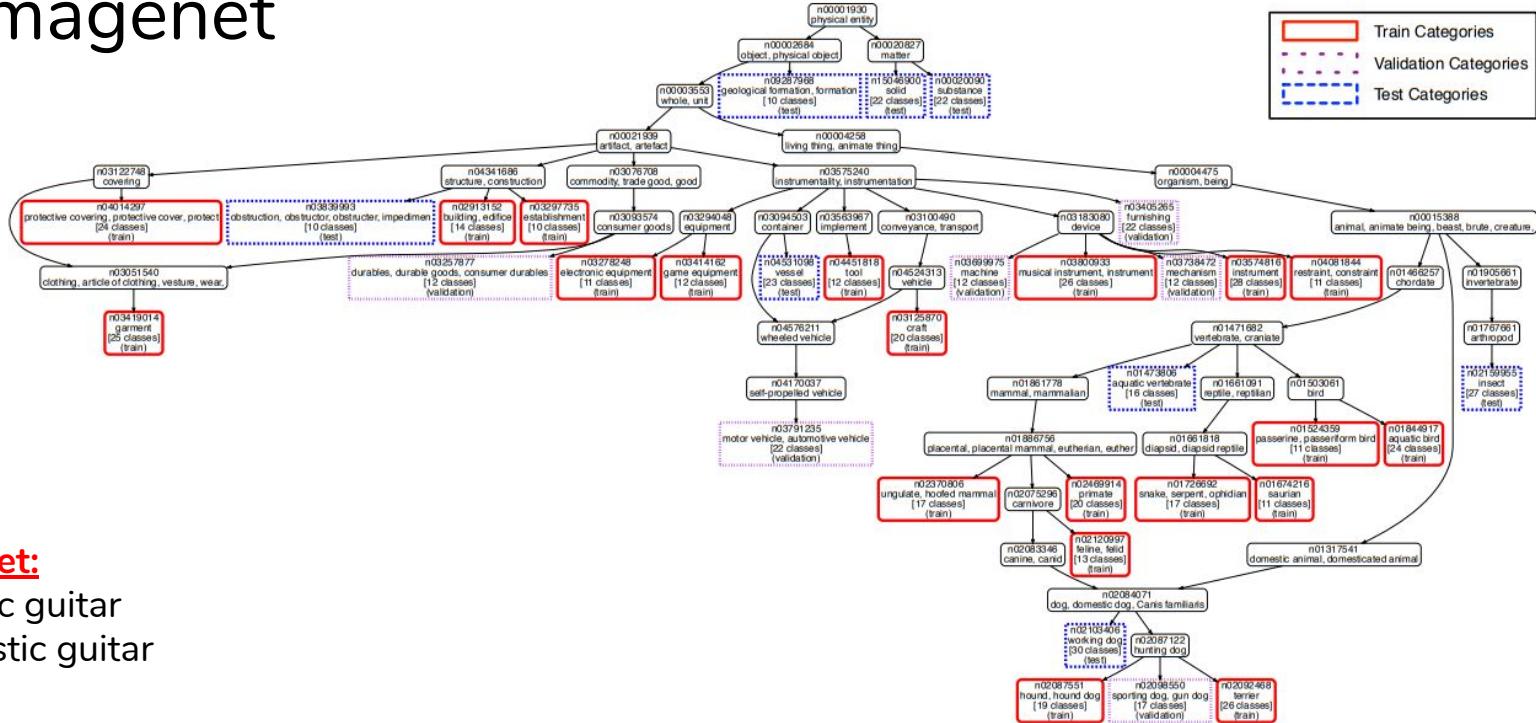
tieredImageNet

**Train Categories:** n02087551 (hound, hound dog), n02092468 (terrier), n02120997 (feline, felid), n02370806 (ungulate, hoofed mammal), n02469914 (primate), n01726692 (snake, serpent, ophidian), n01674216 (saurian), n01524359 (passerine, passeriform bird), n01844917 (aquatic bird), n04081844 (restraint, constraint), n03574816 (instrument), n03800933 (musical instrument, instrument), n03125870 (craft), n04451818 (tool), n03414162 (game equipment), n03278248 (electronic equipment), n03419014 (garment), n03297735 (establishment), n02913152 (building, edifice), n04014297 (protective covering, protective cover, protection).

**Validation Categories:** n02098550 (sporting dog, gun dog), n03257877 (durables, durable goods, consumer durables), n03405265 (furnishing), n03699975 (machine), n03738472 (mechanism), n03791235 (motor vehicle, automotive vehicle),

**Test Categories:** n02103406 (working dog), n01473806 (aquatic vertebrate), n02159955 (insect), n04531098 (vessel), n03839993 (obstruction, obstructor, obstructor, impediment, impedimenta), n09287968 (geological formation, formation), n00020090 (substance), n15046900 (solid).

# tieredImagenet



## minilmageNet:

**Test** - electric guitar  
**Train** - acoustic guitar

## tieredImageNet:

**Test** - musical instruments  
**Train** - farming equipment

# Datasets

- Omniglot
- minilmageNet (600 images from 100 classes)
- tieredImageNet (34 broad categories, each containing 10 to 30 classes)

10% goes to labeled splits

90% goes to unlabelled classes and distractors\*

\*40/60 for minilmageNet

# Datasets

- Omniglot
- minilmageNet (600 images from 100 classes)
- tieredImageNet (34 broad categories, each containing 10 to 30 classes)

10% goes to labeled splits

Much less labelled data than standard few-shot approaches!!!

90% goes to unlabelled classes and distractors\*

\*40/60 for minilmageNet

# Datasets

**N:** Classes

**K:** Labelled samples from each class

**M:** Unlabelled samples from N classes

**H:** Distractors (Unlabelled sample from classes other than N)

**H=N=5**

**M=5** for training & **M=20** for testing

# Baseline Models

Models	Acc.	Acc. w/ D
1. Supervised	$94.62 \pm 0.09$	$94.62 \pm 0.09$
Semi-Supervised Inference	$97.45 \pm 0.05$	$95.08 \pm 0.09$

1. Vanilla Protonet

# Baseline Models

Models	Acc.	Acc. w/ D
1. Supervised	$94.62 \pm 0.09$	$94.62 \pm 0.09$
2. Semi-Supervised Inference	$97.45 \pm 0.05$	$95.08 \pm 0.09$

1. Vanilla Protonet
2. Vanilla Protonet + one step of Soft k-means refinement at  
**test only** (supervised embedding)

# Results: Omniglot

Models	Acc.	Acc. w/ D
Supervised	$94.62 \pm 0.09$	$94.62 \pm 0.09$
Semi-Supervised Inference	$97.45 \pm 0.05$	$95.08 \pm 0.09$
Soft $k$ -Means	$97.25 \pm 0.10$	$95.01 \pm 0.09$
Soft $k$ -Means+Cluster	<b><math>97.68 \pm 0.07</math></b>	$97.17 \pm 0.04$
Masked Soft $k$ -Means	$97.52 \pm 0.07$	<b><math>97.30 \pm 0.08</math></b>

# Results: minilmageNet

Models	1-shot Acc.	5-shot Acc.	1-shot Acc w/ D	5-shot Acc. w/ D
Supervised	$43.61 \pm 0.27$	$59.08 \pm 0.22$	$43.61 \pm 0.27$	$59.08 \pm 0.22$
Semi-Supervised Inference	$48.98 \pm 0.34$	$63.77 \pm 0.20$	$47.42 \pm 0.33$	$62.62 \pm 0.24$
Soft $k$ -Means	<b><math>50.09 \pm 0.45</math></b>	<b><math>64.59 \pm 0.28</math></b>	<b><math>48.70 \pm 0.32</math></b>	<b><math>63.55 \pm 0.28</math></b>
Soft $k$ -Means+Cluster	$49.03 \pm 0.24$	$63.08 \pm 0.18$	<b><math>48.86 \pm 0.32</math></b>	$61.27 \pm 0.24$
Masked Soft $k$ -Means	<b><math>50.41 \pm 0.31</math></b>	<b><math>64.39 \pm 0.24</math></b>	<b><math>49.04 \pm 0.31</math></b>	$62.96 \pm 0.14$

**Table 2:** *miniImageNet* 1/5-shot classification results.

# Results: tieredImageNet

Models	1-shot Acc.	5-shot Acc.	1-shot Acc. w/ D	5-shot Acc. w/ D
Supervised	$46.52 \pm 0.52$	$66.15 \pm 0.22$	$46.52 \pm 0.52$	$66.15 \pm 0.22$
Semi-Supervised Inference	$50.74 \pm 0.75$	$69.37 \pm 0.26$	$48.67 \pm 0.60$	$67.46 \pm 0.24$
Soft $k$ -Means	$51.52 \pm 0.36$	<b><math>70.25 \pm 0.31</math></b>	$49.88 \pm 0.52$	$68.32 \pm 0.22$
Soft $k$ -Means+Cluster	<b><math>51.85 \pm 0.25</math></b>	$69.42 \pm 0.17$	<b><math>51.36 \pm 0.31</math></b>	$67.56 \pm 0.10$
Masked Soft $k$ -Means	<b><math>52.39 \pm 0.44</math></b>	<b><math>69.88 \pm 0.20</math></b>	<b><math>51.38 \pm 0.38</math></b>	<b><math>69.08 \pm 0.25</math></b>

**Table 3:** *tieredImageNet* 1/5-shot classification results.

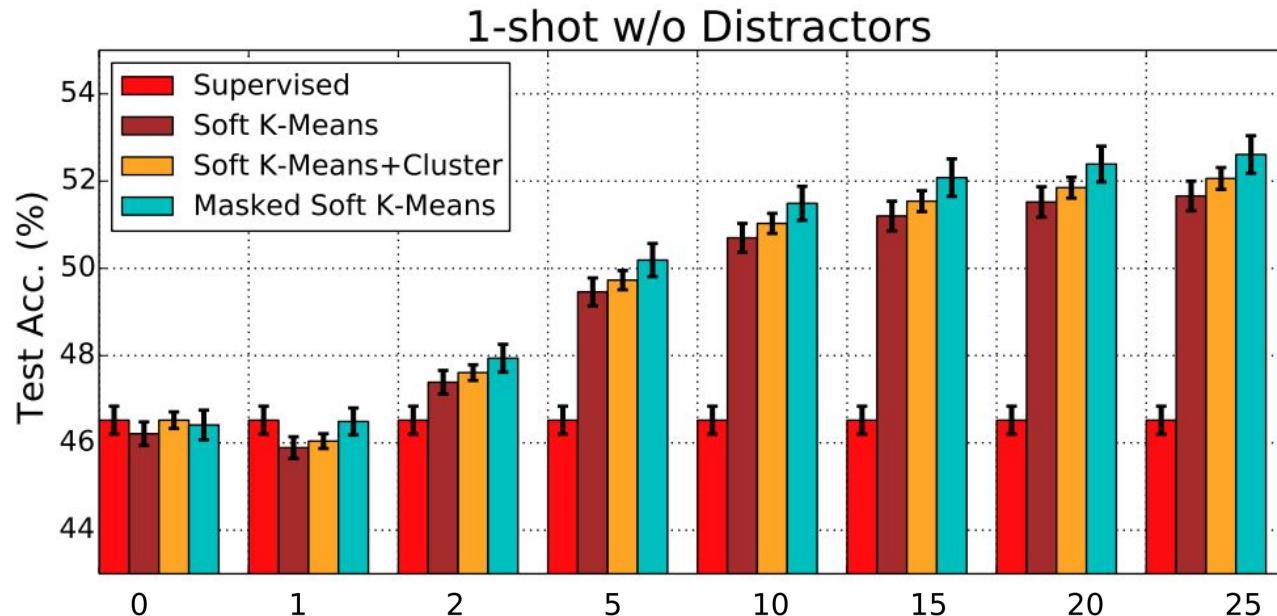
# Results: Other Baselines

Models	Omniglot	<i>miniImageNet</i>		<i>tieredImageNet</i>	
	1-shot	1-shot	5-shot	1-shot	5-shot
1-NN Pixel	$40.39 \pm 0.36$	$26.74 \pm 0.48$	$31.43 \pm 0.51$	$26.55 \pm 0.50$	$30.79 \pm 0.53$
1-NN CNN rnd	$59.55 \pm 0.46$	$24.03 \pm 0.38$	$27.54 \pm 0.42$	$25.49 \pm 0.45$	$30.01 \pm 0.47$
1-NN CNN pre	$52.53 \pm 0.51$	$32.90 \pm 0.58$	$40.79 \pm 0.76$	$32.76 \pm 0.66$	$40.26 \pm 0.67$
LR Pixel	$49.15 \pm 0.39$	$24.50 \pm 0.41$	$33.33 \pm 0.68$	$25.70 \pm 0.46$	$36.30 \pm 0.62$
LR CNN rnd	$57.80 \pm 0.45$	$24.10 \pm 0.50$	$28.40 \pm 0.42$	$26.55 \pm 0.48$	$32.51 \pm 0.52$
LR CNN pre	$48.49 \pm 0.47$	$30.28 \pm 0.54$	$40.27 \pm 0.59$	$34.52 \pm 0.68$	$43.58 \pm 0.72$
ProtoNet	<b><math>94.62 \pm 0.09</math></b>	<b><math>43.61 \pm 0.27</math></b>	<b><math>59.08 \pm 0.22</math></b>	<b><math>46.52 \pm 0.32</math></b>	<b><math>66.15 \pm 0.34</math></b>

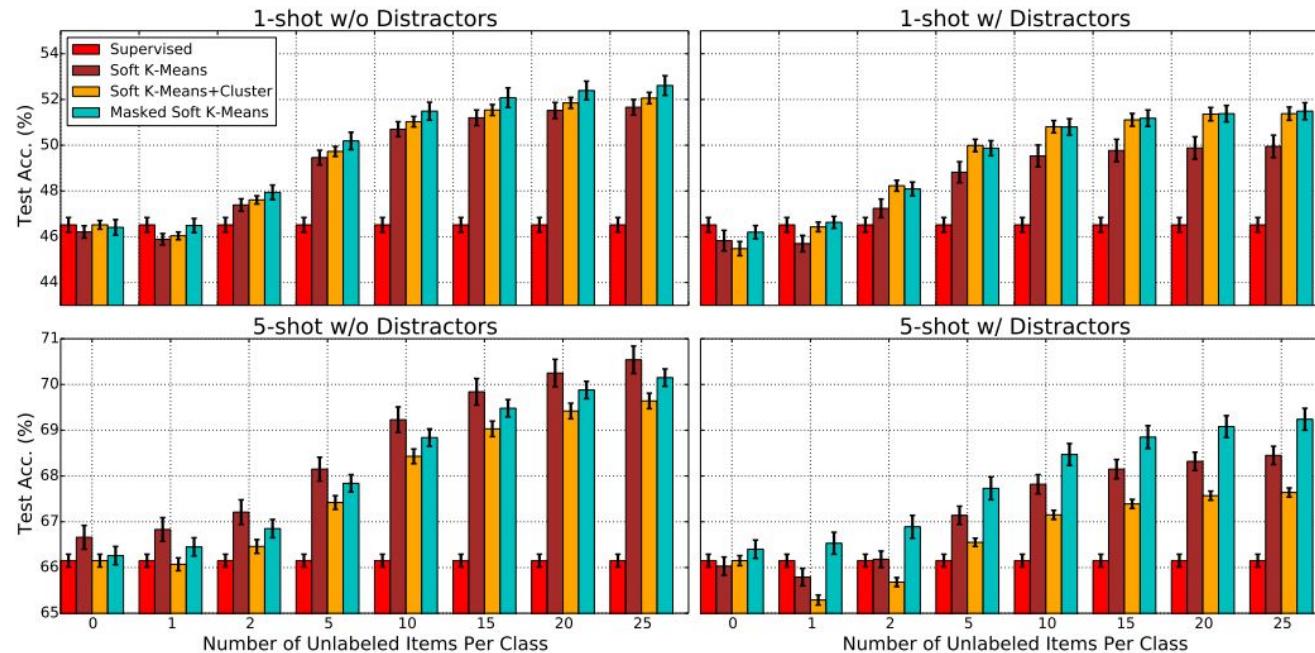
# Results

Models trained with M=5

During meta-test: vary amount of unlabelled examples



# Results



# Conclusions:

1. Achieve state of the art performance over logical baselines on 3 datasets

# Conclusions:

1. Achieve state of the art performance over logical baselines on 3 datasets
2. K-means Masked models perform best with distractors

# Conclusions:

1. Achieve state of the art performance over logical baselines on 3 datasets
2. K-means Masked models perform best with distractors
3. Novel: models extrapolate to increases in amount of labelled data

# Conclusions:

1. Achieve state of the art performance over logical baselines on 3 datasets
2. K-means Masked models perform best with distractors
3. Novel: models extrapolate to increases in amount of labelled data
4. New dataset: tieredImageNet

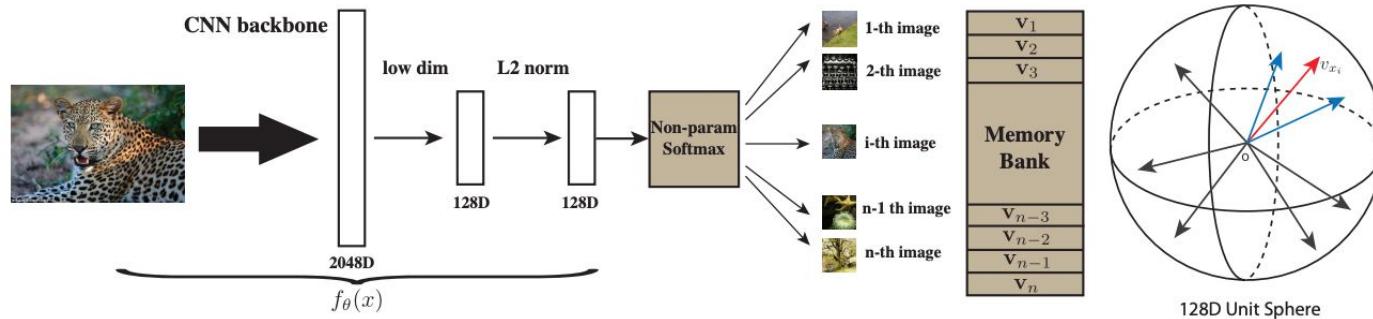
## Critiques:

1. Results are convincing, but the work is actually a relatively straightforward application of (a) **Protonets** and (b) **k-means clustering**
2. **Model Choice:** protonets are very simple. It's not clear what they gained by the simple inductive bias
3. Presented approach does not generalize well beyond classification problems

# Future directions: extension to unsupervised learning

I would be really interested in withholding labels altogether  
Can the model learn how many classes there are?

... and correctly classify them?



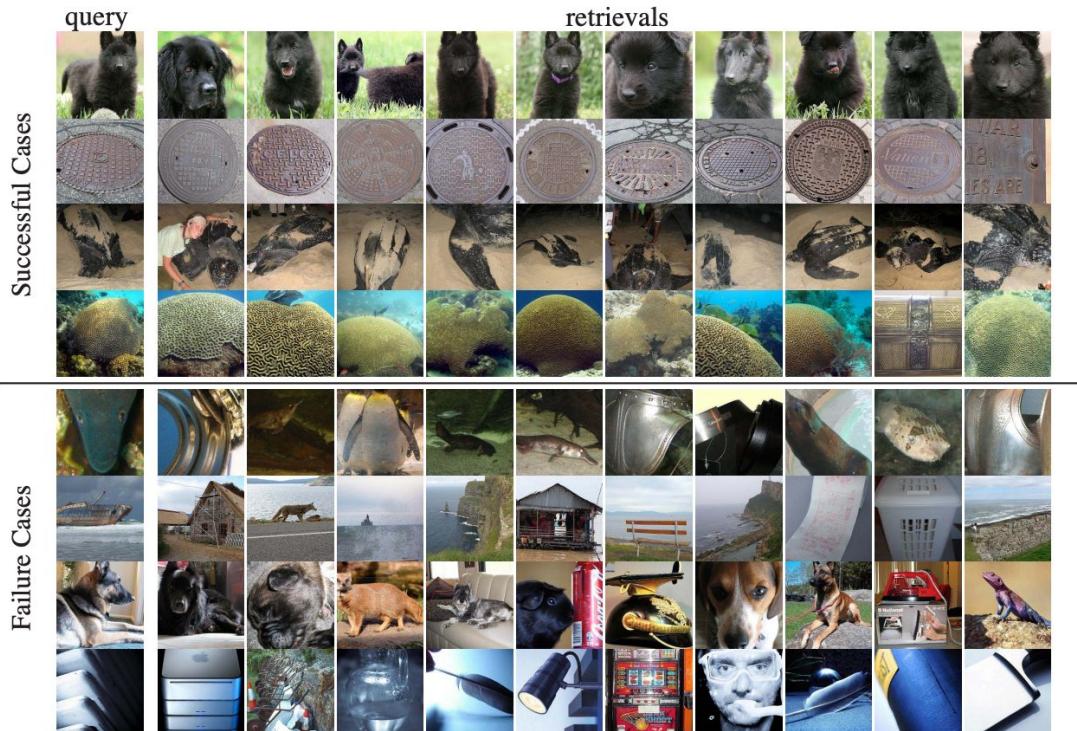
Unsupervised Feature Learning via Non-Parametric Instance Discrimination

Zhirong Wu<sup>\*†</sup>  
\*UC Berkeley / ICSI

Yuanjun Xiong<sup>‡‡</sup>  
†Chinese University of Hong Kong

Stella X. Yu<sup>\*</sup>  
Dahua Lin<sup>†</sup>  
‡Amazon Rekognition

# Future directions: extension to unsupervised learning



Unsupervised Feature Learning via Non-Parametric Instance Discrimination

Zhirong Wu<sup>\*†</sup>  
\*UC Berkeley / ICSI

Yuanjun Xiong<sup>†‡</sup>  
†Chinese University of Hong Kong

Stella X. Yu<sup>\*</sup>  
Dahua Lin<sup>†</sup>  
†Amazon Rekognition

Thank you!

# Supplemental: Accounting for Intra-Cluster Distance

$$\tilde{z}_{j,c} = \frac{\exp\left(-\frac{1}{r_c^2} \|\tilde{\mathbf{x}}_j - \mathbf{p}_c\|_2^2 - A(r_c)\right)}{\sum_{c'} \exp\left(-\frac{1}{r_c^2} \|\tilde{\mathbf{x}}_j - \mathbf{p}_{c'}\|_2^2 - A(r_{c'})\right)}, \text{ where } A(r) = \frac{1}{2} \log(2\pi) + \log(r) \quad (6)$$