

# Review of “Data-Efficient Hierarchical Reinforcement Learning”

28th October, 2019

# Outlines

- Motivation
- Method Overview
- Experimental Results
- Q&A and Discussion

# Motivation

## What is HRL?

Hierarchical Reinforcement Learning will train multiple layers of policies, in which the policy at higher layer generates small goals to policy at lower layer. In particular, the policy at the lowest layer directly interacts with the environment.

# Motivation

## What is HRL?

Hierarchical Reinforcement Learning will train multiple layers of policies, in which the policy at higher layer generates small goals to policy at lower layer. In particular, the policy at the lowest layer directly interacts with the environment.

## Problems of Current HRL Algorithms

- ① They are mostly on-policy algorithms, which are data-inefficient.
- ② They mostly need a lot of task-specific designs, which are not general.

# Motivation

## What is HRL?

Hierarchical Reinforcement Learning will train multiple layers of policies, in which the policy at higher layer generates small goals to policy at lower layer. In particular, the policy at the lowest layer directly interacts with the environment.

## Problems of Current HRL Algorithms

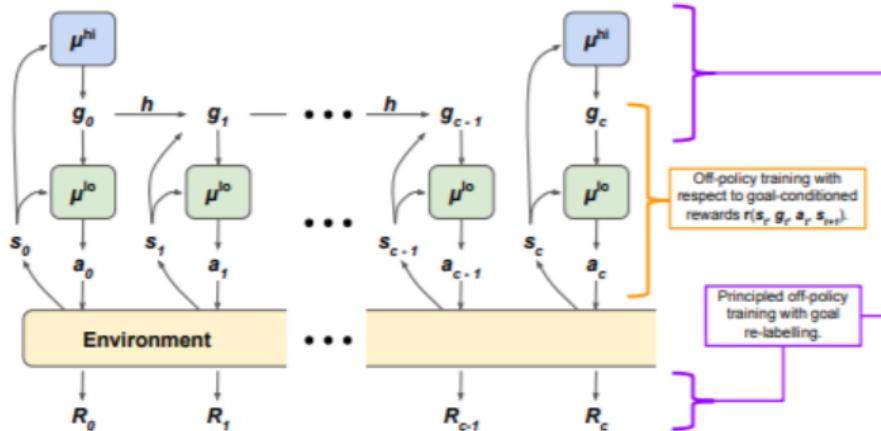
- ① They are mostly on-policy algorithms, which are data-inefficient.
- ② They mostly need a lot of task-specific designs, which are not general.

## Challenge for Designing Off-policy HRL Algorithm

- ① Past high-level goals do not correspond to current low-level agent.
- ② Training can be unstable.

# Model and Method Overview

A two-layer HRL algorithm.



1. Collect experience  $s_t, g_t, a_t, R_t, \dots$
2. Train  $\mu^{lo}$  with experience transitions  $(s_t, g_t, a_t, r_t, s_{t+1}, g_{t+1})$  using  $g_t$  as additional state observation and reward given by goal-conditioned function  $r_t = r(s_t, g_t, a_t, s_{t+1}) = -\|s_t + g_t - s_{t+1}\|_2$ .
3. Train  $\mu^{hi}$  on temporally-extended experience  $(s_t, \tilde{g}_t, \sum R_{t:t+c-1}, s_{t+c})$ , where  $\tilde{g}_t$  is re-labelled high-level action to maximize probability of past low-level actions  $a_{t:t+c-1}$ .
4. Repeat.

## Model and Method Overview (Continued)

The training for high-level and low-level policies can be seen as training on two different but related MDPs.

### High-level MDP

Data buffer content:  $(s_t, \tilde{g}_t, \sum R_{t:t+c-1}, s_{t+c})$

- $s_t$ : MDP state, which is the same as environment state
- $s_{t+c}$ : state after transition, which is environment state after  $c$  steps of transition
- $g_t$ : action, functioning as the goal of low-level policy

$$g_t = \begin{cases} \mu^{high}(s_t) & \text{if } t \equiv 0 \pmod c \\ s_{t-1} + g_{t-1} - s_t & \text{otherwise} \end{cases}$$

- $\tilde{g}_t$ : action used for high-level training, modified by off-policy correction (more details later)
- $\sum R_{t:t+c-1}$ : MDP reward, summation of  $c$  steps' environment reward.

# Model and Method Overview (Continued)

## Low-level MDP

Data Buffer Content:  $((s_t, g_t), a_t, r_t, (s_{t+1}, g_{t+1}))$

- $(s_t, g_t)$ : MDP state, where  $s_t$  is environment state and  $g_t$  is the action of high-level policy
- $a_t$ : action, which interacts directly with the environment
- $r_t$ : MDP reward, which is goal-conditioned, defined as  $r_t = -\|s_t + g_t - s_{t+1}\|_2$
- $(s_{t+1}, g_{t+1})$ : state after transition

# Model and Method Overview (Continued)

## Low-level MDP

Data Buffer Content:  $((s_t, g_t), a_t, r_t, (s_{t+1}, g_{t+1}))$

- $(s_t, g_t)$ : MDP state, where  $s_t$  is environment state and  $g_t$  is the action of high-level policy
- $a_t$ : action, which interacts directly with the environment
- $r_t$ : MDP reward, which is goal-conditioned, defined as  $r_t = -\|s_t + g_t - s_{t+1}\|_2$
- $(s_{t+1}, g_{t+1})$ : state after transition

## Instantiation Algorithm

Separately train high-level and low-level policies by using an off-policy algorithm. Specifically, the authors used TD3 algorithm, which is a variant of DDPG.

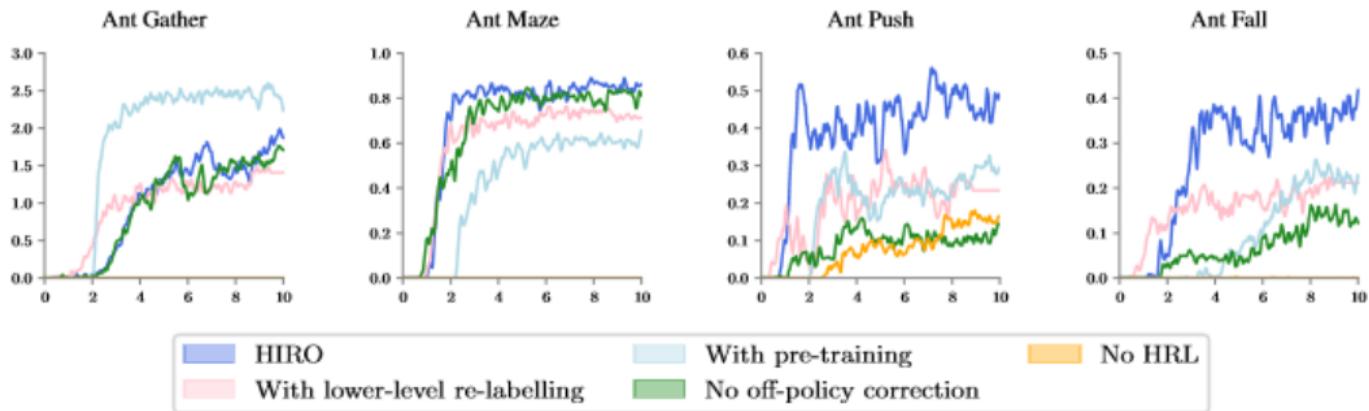
## Off-policy Correction

As the lower-level policy evolves, the lower-level actions taken to achieve the same goal varies, which in turns results in different states and rewards collected along the way. This invalidates old high-level transitions  $(s_t, g_t, \sum R_{t:t+c-1}, s_{t+c})$ .

- Need to relabel the goal for high-level training
- Choose  $\tilde{g}_t$  such that
  - ▶ The probability  $\mu^{\text{low}}(a_{t:t+c-1}|s_{t:t+c-1}, \tilde{g}_{t:t+c-1})$  is maximized
  - ▶ The intermediate goals are calculated by  $\tilde{g}_{t+j+1} = s_{t+j} + \tilde{g}_{t+j} - s_{t+j+1}$ ,  $\forall j : 0 \leq j \leq c - 2$
- Practically, in authors' implementation, the maximization is taken over 10 candidate goals, in which eight are sampled Gaussian centered at  $s_{t+c} - s_t$ , one is the original  $g_t$  and the last one is  $s_{t+c} - s_t$

[Video Link](#)

# Comparison with variants of the HIRO model



## Comparison with other models

- Comparison with FuN, SNN4HRL VIME, on rewards of the best policy obtained in 10M steps of training, averaged over 10 randomly seeded trials with standard error

	<b>Ant Gather</b>	<b>Ant Maze</b>	<b>Ant Push</b>	<b>Ant Fall</b>
HIRO	<b>3.02±1.49</b>	<b>0.99±0.01</b>	<b>0.92±0.04</b>	<b>0.66±0.07</b>
FuN representation	0.03 ± 0.01	0.0 ± 0.0	0.0 ± 0.0	0.0 ± 0.0
FuN transition PG	0.41 ± 0.06	0.0 ± 0.0	0.56 ± 0.39	0.01 ± 0.02
FuN cos similarity	0.85 ± 1.17	0.16 ± 0.33	0.06 ± 0.17	0.07 ± 0.22
FuN	0.01 ± 0.01	0.0 ± 0.0	0.0 ± 0.0	0.0 ± 0.0
SNN4HRL	1.92 ± 0.52	0.0 ± 0.0	0.02 ± 0.01	0.0 ± 0.0
VIME	1.42 ± 0.90	0.0 ± 0.0	0.02 ± 0.02	0.0 ± 0.0

# Summary and Critiques

## Summary of Strengths

- This paper propose a novel way to do off-policy HRL high-level training, by correcting goals used in high-level training.
- Data-Efficient -off-policy
- Generality -goals are learned and proposed automatically by the higher-level controllers

# Summary and Critiques

## Summary of Strengths

- This paper propose a novel way to do off-policy HRL high-level training, by correcting goals used in high-level training.
- Data-Efficient -off-policy
- Generality -goals are learned and proposed automatically by the higher-level controllers

## Critiques of Weakness

- In the Ant Gather case, with the high uncertainty, it is less sure that HIRO performs much better.
- Would be more convincing with more experiments on other settings
- Lack of discussion on the aspects of the models in the experiment section
- Discussion on the settings where this model is particularly helpful and where it might fail or not behave as expected

# Discussion and Q&A

# **Thank you very much!**